

# Code

---

## FlinkKafkaProducer.java

```

package wikiedits;

import org.apache.flink.api.common.functions.FilterFunction;
import org.apache.flink.api.common.functions.MapFunction;
import org.apache.flink.api.common.functions.ReduceFunction;
import org.apache.flink.api.common.serialization.SimpleStringSchema;
import org.apache.flink.api.java.tuple.Tuple2;
import org.apache.flink.streaming.api.datastream.DataStream;
import
org.apache.flink.streaming.api.environment.StreamExecutionEnvironment;
import
org.apache.flink.streaming.api.windowing.assigners.TumblingEventTimeWindows;
import org.apache.flink.streaming.api.windowing.time.Time;
import org.apache.flink.streaming.connectors.kafka.FlinkKafkaProducer011;
import
org.apache.flink.streaming.connectors.kafka.partitioner.FlinkKafkaPartitioner;
import org.apache.flink.streaming.connectors.wikiedits.WikipediaEditEvent;
import
org.apache.flink.streaming.connectors.wikiedits.WikipediaEditsSource;
import org.apache.flink.util.Preconditions;

import java.util.Optional;

import static
org.apache.flink.streaming.connectors.kafka.FlinkKafkaProducerBase.getPropertiesFromBrokerList;

public class FlinkKafkaProducer {

    public static void main(String[] args) throws Exception {
        // set up the streaming execution environment
        final StreamExecutionEnvironment env =
StreamExecutionEnvironment.getExecutionEnvironment();
        DataStream<WikipediaEditEvent> edits = env.addSource(new
WikipediaEditsSource());

        DataStream<String> stream =
edits.map(WikipediaEditEvent::toString);
        FlinkKafkaProducer011<String> myProducer = new
FlinkKafkaProducer011<String>(
            "wiki-edits", // target topic
            new SimpleStringSchema(), // serialization schema
            getPropertiesFromBrokerList("localhost:9092"), // broker
list

```

```

        Optional.of(new CustomPartitioner<>()); // round robin
partitioner
        stream.addSink(myProducer);

        DataStream<Tuple2<String, Integer>> result = edits
            // project the event user and the diff
            .map(new MapFunction<WikipediaEditEvent, Tuple2<String,
Integer>>() {
                @Override
                public Tuple2<String, Integer> map(WikipediaEditEvent
event) {
                    return new Tuple2<>(
                        event.getUser(), event.getByteDiff());
                }
            })
            // group by user
            .keyBy(0)
            // aggregate changes per user
            .reduce(new ReduceFunction<Tuple2<String, Integer>>() {
                @Override
                public Tuple2<String, Integer> reduce(Tuple2<String,
Integer> e1, Tuple2<String, Integer> e2) {
                    return new Tuple2<>(e1.f0, e1.f1 + e2.f1);
                }
            })
            // filter out negative byte changes
            .filter(new FilterFunction<Tuple2<String, Integer>>() {
                @Override
                public boolean filter(Tuple2<String, Integer> e)
throws Exception {
                    return e.f1 >= 0;
                }
            });

        // execute program
        env.execute("Flink Streaming Java API Skeleton");
    }

    static class CustomPartitioner<T> extends FlinkKafkaPartitioner<T> {
        private int next = 0;

        @Override
        public int partition(T record, byte[] key, byte[] value, String
targetTopic, int[] partitions) {
            Preconditions.checkArgument(partitions != null &&
partitions.length > 0, "Partitions of the target topic is empty.");
            this.next = (this.next + 1) % partitions.length;
            return partitions[this.next];
        }
    }
}

```

## FlinkKafkaConsumer.java

```

package wikiedits;

import java.util.Properties;
import java.util.regex.Matcher;
import java.util.regex.Pattern;

import org.apache.flink.api.common.functions.FilterFunction;
import org.apache.flink.api.common.functions.MapFunction;
import org.apache.flink.api.common.functions.ReduceFunction;
import
org.apache.flink.api.common.serialization.AbstractDeserializationSchema;
import org.apache.flink.api.java.tuple.Tuple2;
import org.apache.flink.streaming.api.datastream.DataStream;
import
org.apache.flink.streaming.api.environment.StreamExecutionEnvironment;
import org.apache.flink.streaming.api.windowing.time.Time;
import org.apache.flink.streaming.connectors.kafka.FlinkKafkaConsumer011;
import org.apache.flink.streaming.connectors.wikiedits.WikipediaEditEvent;

public class FlinkKafkaConsumer {

    public static void main(String[] args) throws Exception {
        // set up the streaming execution environment
        final StreamExecutionEnvironment env =
StreamExecutionEnvironment.getExecutionEnvironment();

        Properties kafkaProps = new Properties();
        kafkaProps.setProperty("zookeeper.connect", "localhost:2181");
        kafkaProps.setProperty("bootstrap.servers", "localhost:9092");
        kafkaProps.setProperty("group.id", "test-consumer-group");
        // always read the Kafka topic from the start
        kafkaProps.setProperty("auto.offset.reset", "earliest");
        DataStream<WikipediaEditEvent> edits = env
            .addSource(new FlinkKafkaConsumer011<>("wiki-edits",
                new CustomDeserializationSchema(), kafkaProps));

        DataStream<Tuple2<String, Integer>> result = edits
            // project the event user and the diff
            .map(new MapFunction<WikipediaEditEvent, Tuple2<String,
Integer>>() {
                @Override
                public Tuple2<String, Integer> map(WikipediaEditEvent
event) {
                    return new Tuple2<>(
                        event.getUser(), event.getByteDiff());
                }
            })
            // group by user
            .keyBy(0)

```

```

        // tumbling event-time windows
        .timeWindow(Time.seconds(10))
        // aggregate changes per user
        .reduce(new ReduceFunction

```

```
        return new WikipediaEditEvent(Long.valueOf(m.group(1)),
m.group(2), m.group(3), m.group(4), m.group(5),
Integer.valueOf(m.group(6)), m.group(7), isMinor, isNew, isUnpatrolled,
isBotEdit, isSpecial, isTalk);
    } else {
        throw new RuntimeException("error in parsing!");
    }
}
}
```